# Haptica:
# Own Your Context

# Context is the bottleneck.

- Status updates, handoffs, and re-explanations eat hours every week.
- AI agents need context pasted in every session.
- Progress is invisible until you narrate it.

# Haptica is a context engine for your work.

- Dynamic memory, built from your screen.
- A live wiki for every project, updating as you work.
- Coordinate teams and multi-agent workflows.

CAPTURE

### Your screen

browser tabs
terminal
editor
chat

→

PROCESS

### Haptica agent

OCR
chunking
indexing
synthesis

→

ORGANIZE

### Project wikis

per project
timeline
decisions
live

→

SERVE

### Humans + agents

share links
via MCP
Slack + Linear
auth-gated

# One project Haptica tracked.

1Context  Search pages, books, tags...  Shared · 5  Light

CONTENTS


Mission
Architecture split
- Relay canonical
- App canonical
- GCD future canonical

HTTP routes
SSE streaming contract
Warm sessions
Deployment topology
Service management
Verification
What the relay is NOT
Related pages


TEAM-DOCS › GUARDIAN › **LAB-RELAY**

## Guardian Lab Relay

*A flexible execution node for Guardian — not a second product backend.*

The **Guardian Lab Relay** is a thin HTTP/SSE relay that runs Claude Code sessions from the Guardian app against the live Guardian project environment on the B650 host. It pins `cwd` to the active Guardian project so MCP and config are present, persists conversations and turns to SQLite, and shapes raw agent behavior into product-friendly streaming events for the app.

It is intentionally small. It is *not* meant to become the long-term canonical cloud backend for the product — that role is reserved for GCD.

### Fast facts

| | |
|---|---|
| STATUS | Live on B650 host |
| REPO | `hapticasensorics/guardian-lab-relay` |
| STACK | Node.js + Express + SQLite + SSE |
| LIVE MODEL | `claude-opus-4-6[1m]` fast mode enabled |
| SLOTS | 3 (blue / green / experimental) |
| PORTS | 3000 · 3001 · 3002 · 3010 (router) |
| TAILNET HTTPS | `haptica-b650-pro-rs.tail5714b3.ts.net` |
| LAST HARNESS | 55 passed · 0 failed |
| MAINTAINER | paul@haptica.ai |

### Mission

The relay should be the best place to:

- Run Claude/agent sessions against the live Guardian environment.
- Prototype app-to-agent interaction models quickly.
- Experiment with warm-session and memory-loading strategies.
- Shape streaming behavior into product-friendly UX contracts.
- Orchestrate Guardian tools and MCP integrations.
- Keep relay continuity state durable and testable.
- Support trusted-user demos, research, and fast iteration.

> *The lab relay is a stepping stone, not a dead end. We want to learn fast locally without creating migration debt everywhere.*
>
> — lab-relay-mission.md

# Haptica already knows.

```
$ help me debug this

Sure, can you tell me what you're
working on, what language, and
paste the error?
```

without haptica

```
$ help me debug this

You've been working on the cross-device
alignment pipeline since Tuesday. Last
night you decided to keep Whisper+Gemini
reversible. The failing test is in
scene_claims.py:142 — same module you
refactored yesterday.

> haptica context: 4.2h session · 3 tools · 2 decisions tracked
```

with haptica

# Market size.

*The AI coding tools category went $0 to $7B+ in three years.*

**TODAY**

Cursor · 1M paying · $2B ARR

Claude Code · $2.5B ARR · **$20 / seat**

Codex · 3M+ weekly active users

Copilot Business · 1.8M+ seats

**OUR WEDGE**

$1.2B

5M+ devs running AI coding agents · at $20 / seat

**Next:** every knowledge worker with an AI agent. $24B+ TAM by 2028.

# Business model

**$20** *per seat, per month*

3-seat minimum.

- Dev-first. Initial motion: founder-to-founder outreach.
- Built for devs. Sold to teams.

CURRENT PILOT pandatron 30-day pilot, engineering team

# Start with devs running agentic coding agents.

WEDGE

## 100–500

Engineers running Claude Code, Cursor, and Codex inside company teams.

BRIDGE

## 500–5,000

High-context knowledge workers

BEYOND

## 5,000+

Phone + mic: full-context capture across every device

# Why now

01 AI is very capable.

02 Users are delegating real work to AI agents.

03 Context is the limiting factor.

# How we win.

*Most AI products ask you to change your workflow. Ours asks for nothing.*

- One command to install.
- Works through the tools you already use: Claude Code, Cursor, ChatGPT. MCP-native.
- Context compounds daily.

# Haptica Team

Paul Han
Founder

Former datacenter hardware engineer at Meta. Nuclear weapons programs at Lawrence Livermore National Laboratory. Yale, electrical engineering and political science.

Jackie Oliver
Co-Founder

Former data engineer building production ML scoring models and pipelines for fintech credit-risk products. UVA, economics and philosophy.

THE ASK

# $1M *$15M cap · Standard YC SAFE*

**SIGN SAFE →**

haptica.ai/safe

**USE OF FUNDS**

- First technical hire: senior engineer
- Inference & infrastructure (OCR, wiki synthesis, MCP hosting)
- Developer community, pilots, outreach
- Product (new integrations, sharing, Windows)

**MILESTONES**

- 100 daily active users
- Published context benchmark
- $1M ARR

# *Thank you.*

paul@haptica.ai | jackieoliver@haptica.ai