Haptica:
Own Your Context

Context is the bottleneck.

- Status updates, handoffs, and re-explanations eat hours every week.
- AI agents need context pasted in every session.
- Progress is invisible until you narrate it.

Haptica is a context engine for your work.

- Dynamic memory, built from your screen.
- A live wiki for every project, updating as you work.
- Coordinate teams and multi-agent workflows.

CAPTURE

Your screen

browser tabs
terminal
editor
chat

→

PROCESS

Haptica agent

OCR
chunking
indexing
synthesis

→

ORGANIZE

Project wikis

per project
timeline
decisions
live

→

SERVE

Humans + agents

share links
via MCP
Slack + Linear
auth-gated

One project Haptica tracked.

1Context Search pages, books, tags... Shared · 5 Light

CONTENTS

Mission
Architecture split
- Relay canonical
- App canonical
- GCD future canonical

HTTP routes
SSE streaming contract
Warm sessions
Deployment topology
Service management
Verification
What the relay is NOT
Related pages

TEAM-DOCS › GUARDIAN › **LAB-RELAY**

Guardian Lab Relay

A flexible execution node for Guardian — not a second product backend.

The **Guardian Lab Relay** is a thin HTTP/SSE relay that runs Claude Code sessions from the Guardian app against the live Guardian project environment on the B650 host. It pins `cwd` to the active Guardian project so MCP and config are present, persists conversations and turns to SQLite, and shapes raw agent behavior into product-friendly streaming events for the app.

It is intentionally small. It is *not* meant to become the long-term canonical cloud backend for the product — that role is reserved for GCD.

Fast facts

STATUS	Live on B650 host
REPO	`hapticasensorics/guardian-lab-relay`
STACK	Node.js + Express + SQLite + SSE
LIVE MODEL	`claude-opus-4-6[1m]` fast mode enabled
SLOTS	3 (blue / green / experimental)
PORTS	3000 · 3001 · 3002 · 3010 (router)
TAILNET HTTPS	`haptica-b650-pro-rs.tail5714b3.ts.net`
LAST HARNESS	55 passed · 0 failed
MAINTAINER	paul@haptica.ai

Mission

The relay should be the best place to:

- Run Claude/agent sessions against the live Guardian environment.
- Prototype app-to-agent interaction models quickly.
- Experiment with warm-session and memory-loading strategies.
- Shape streaming behavior into product-friendly UX contracts.
- Orchestrate Guardian tools and MCP integrations.
- Keep relay continuity state durable and testable.
- Support trusted-user demos, research, and fast iteration.

> *The lab relay is a stepping stone, not a dead end. We want to learn fast locally without creating migration debt everywhere.*
>
> — lab-relay-mission.md

Haptica already knows.

```
$ help me debug this

Sure, can you tell me what you're
working on, what language, and
paste the error?
```

without haptica

```
$ help me debug this

You've been working on the cross-device
alignment pipeline since Tuesday. Last
night you decided to keep Whisper+Gemini
reversible. The failing test is in
scene_claims.py:142 — same module you
refactored yesterday.

> haptica context: 4.2h session · 3 tools · 2 decisions tracked
```

with haptica

Market size.

The AI coding tools category went $0 to $7B+ in three years.

TODAY

Cursor · 1M paying · $2B ARR

Claude Code · $2.5B ARR · **$20 / seat**

Codex · 3M+ weekly active users

Copilot Business · 1.8M+ seats

OUR WEDGE

$1.2B

5M+ devs running AI coding agents · at $20 / seat

Next: every knowledge worker with an AI agent. $24B+ TAM by 2028.

Business model

$20 *per seat, per month*

3-seat minimum.

- Dev-first. Initial motion: founder-to-founder outreach.
- Built for devs. Sold to teams.

CURRENT PILOT pandatron 30-day pilot, engineering team

Start with devs running agentic coding agents.

WEDGE

100–500

Engineers running Claude Code, Cursor, and Codex inside company teams.

BRIDGE

500–5,000

High-context knowledge workers

BEYOND

5,000+

Phone + mic: full-context capture across every device

Why now

01 AI is very capable.

02 Users are delegating real work to AI agents.

03 Context is the limiting factor.

How we win.

Most AI products ask you to change your workflow. Ours asks for nothing.

- One command to install.
- Works through the tools you already use: Claude Code, Cursor, ChatGPT. MCP-native.
- Context compounds daily.

Haptica Team

Paul Han
Founder

Former datacenter hardware engineer at Meta. Nuclear weapons programs at Lawrence Livermore National Laboratory. Yale, electrical engineering and political science.

Jackie Oliver
Co-Founder

Former data engineer building production ML scoring models and pipelines for fintech credit-risk products. UVA, economics and philosophy.

THE ASK

$1M

$15M cap · Standard YC SAFE

SIGN SAFE →

haptica.ai/safe

USE OF FUNDS

- First technical hire: senior engineer
- Inference & infrastructure (OCR, wiki synthesis, MCP hosting)
- Developer community, pilots, outreach
- Product (new integrations, sharing, Windows)

MILESTONES

- 100 daily active users
- Published context benchmark
- $1M ARR

Thank you.

paul@haptica.ai | jackieoliver@haptica.ai